


Dataset Composition and Evaluation Design

Patient-level split for the paired training experiment and separate external domain-shift evaluation

Verified sample counts



Interpretation boundaries

Training objective	Self-supervised CT reconstruction and feature alignment
Held-out benchmark	Zenodo: n=3, same source domain
External benchmark	IBSR: n=40, different domain
Labels	No verified WHO-grade labels
Registration	Not performed in current DICOM track
Classification metrics	Not reported or claimed